

# **Optimizing Symbolic Execution Through Taint Analysis and Path Prioritization**

Bachelor thesis

Natural Science Faculty of the University of Basel  
Department of Mathematics and Computer Science  
Databases and Information Systems (DBIS) Group  
<https://dbis.dmi.unibas.ch/>

Examiner: Dr. Marco Vogt  
Supervisor: Prof. Dr. Christopher Scherb

Ruben Hutter  
[ruben.hutter@unibas.ch](mailto:ruben.hutter@unibas.ch)  
2020-065-934

02.07.2025

## Acknowledgments

I would like to express my sincere gratitude to Prof. Dr. Christopher Scherb for his supervision and guidance throughout this thesis. His expertise in program analysis and symbolic execution provided essential direction for this research.

I am grateful to Dr. Marco Vogt for his role as examiner and for facilitating the opportunity to pursue this research topic. His feedback and suggestions helped refine both the technical approach and the presentation of this work.

I would like to thank Ivan Giangreco for providing the LaTeX thesis template used for this document, which greatly facilitated the formatting and structure of this work.

I also acknowledge my fellow student Nico Bachmann for developing the Schnauzer visualization library, which enhanced the presentation and analysis of the results in this work.

I thank my family and friends for their encouragement and support during my studies, which made completing this thesis possible.

Finally, I acknowledge the developers of the angr binary analysis framework, whose comprehensive platform enabled the implementation of the techniques presented in this work.

# Abstract

Symbolic execution is a powerful program analysis technique widely used for vulnerability discovery and test case generation. However, its practical application is often hampered by scalability issues, primarily due to the "path explosion problem" where the number of possible execution paths grows exponentially with program complexity. This thesis addresses this fundamental challenge by proposing an optimized approach to symbolic execution that integrates taint analysis and path prioritization.

The core contribution is a novel exploration strategy that moves away from uniform path exploration towards targeted analysis of security-critical program behaviors. The approach prioritizes execution paths originating from memory allocations and user input processing points, as these represent common sources of vulnerabilities. By leveraging dynamic taint analysis, the system identifies and tracks data flow from these critical sources, enabling the symbolic execution engine to focus computational resources on paths influenced by tainted data while deprioritizing paths with no dependency on external inputs.

The implementation integrates this taint-guided exploration strategy with the angr symbolic execution framework, introducing a scoring mechanism that dynamically adjusts path prioritization based on taint propagation. The effectiveness of this optimization is evaluated through comparative analysis, examining runtime efficiency, path coverage quality, and vulnerability discovery capabilities. Results demonstrate that this approach can significantly reduce the search space while maintaining or improving the detection of security-relevant program behaviors, making symbolic execution more practical for large and complex software systems.

# Table of Contents

<b>Acknowledgments</b>	<b>ii</b>
<b>Abstract</b>	<b>iii</b>
<b>1 Introduction</b>	<b>1</b>
<b>2 Background</b>	<b>4</b>
2.1 Symbolic Execution . . . . .	4
2.2 Program Vulnerability Analysis . . . . .	5
2.3 Taint Analysis . . . . .	5
2.4 Control Flow Analysis . . . . .	5
2.5 Angr Framework . . . . .	6
<b>3 Related Work</b>	<b>7</b>
3.1 Optimization Approaches . . . . .	7
3.1.1 State Space Reduction . . . . .	7
3.1.2 Performance and Compositional Analysis . . . . .	7
3.2 Integration Approaches . . . . .	8
3.3 Security-Focused Targeting and Research Gap . . . . .	8
<b>4 Taint-Guided Exploration</b>	<b>9</b>
4.1 Core Approach . . . . .	9
4.2 Taint Source Recognition . . . . .	10
4.3 Dynamic Taint Tracking . . . . .	10
4.4 Path Prioritization . . . . .	11
4.4.1 Adaptive State Pool Management . . . . .	12
4.5 Exploration Depth Control and Vulnerability Probability . . . . .	13
<b>5 Implementation</b>	<b>15</b>
5.1 System Architecture . . . . .	15
5.1.1 Core Design Principles . . . . .	16
5.2 Angr Integration Layer . . . . .	16
5.2.1 Custom Exploration Technique . . . . .	16
5.2.2 State Scoring Implementation . . . . .	17
5.3 Taint Tracking Implementation . . . . .	17

---

5.3.1	Taint Source Detection . . . . .	17
5.3.2	Taint Propagation Engine . . . . .	18
5.3.3	Inter-procedural Taint Flow . . . . .	18
5.4	Configuration and Usage . . . . .	19
5.4.1	Configuration System . . . . .	19
5.4.2	Integration Workflow . . . . .	19
5.4.3	Performance Considerations . . . . .	19
5.5	Implementation Validation . . . . .	20
<b>6</b>	<b>Evaluation</b>	<b>21</b>
6.1	Experimental Design . . . . .	21
6.1.1	Research Questions . . . . .	21
6.1.2	Evaluation Metrics . . . . .	21
6.2	Benchmark Programs . . . . .	21
6.2.1	Synthetic Benchmarks . . . . .	21
6.2.2	Real-World Programs . . . . .	21
6.3	Experimental Results . . . . .	21
6.3.1	Comparison with Standard Symbolic Execution . . . . .	21
6.3.2	Ablation Studies . . . . .	21
6.4	Case Studies . . . . .	21
6.4.1	Buffer Overflow Discovery . . . . .	21
6.4.2	Format String Vulnerability . . . . .	21
<b>7</b>	<b>Conclusion</b>	<b>22</b>
<b>8</b>	<b>Future Work</b>	<b>23</b>
<b>9</b>	<b>Usage of AI</b>	<b>24</b>
	<b>Bibliography</b>	<b>25</b>
	<b>Appendix A Appendix</b>	<b>26</b>

# 1

## Introduction

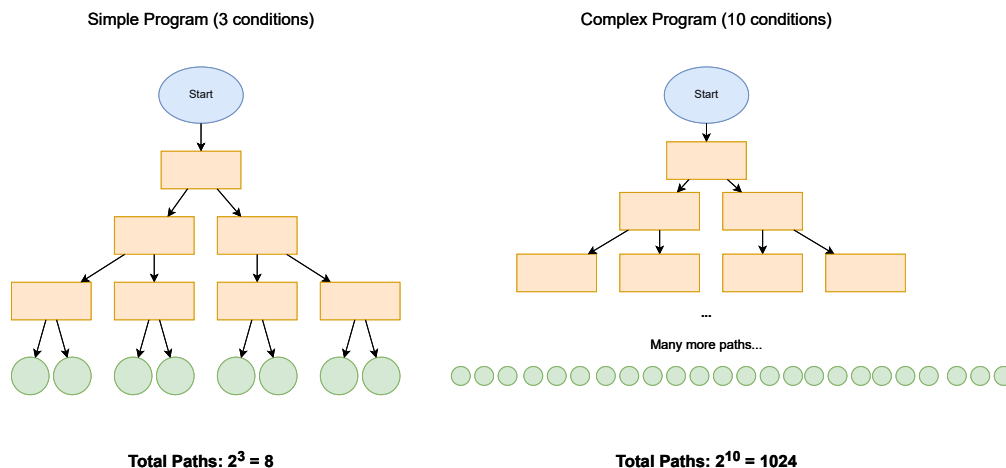
In today's interconnected digital landscape, software security has become a critical concern as applications handle increasingly sensitive data and operate in hostile environments. The discovery of security vulnerabilities before deployment is essential to prevent exploitation by malicious actors, yet traditional testing approaches often fail to comprehensively explore all possible execution scenarios, leaving potential vulnerabilities undiscovered.

Among program analysis techniques, symbolic execution has emerged as a particularly powerful approach for automated vulnerability discovery. Unlike traditional testing that executes programs with concrete input values, symbolic execution treats inputs as mathematical symbols and tracks how these symbols propagate through program computations. When encountering conditional branches, the symbolic execution engine explores multiple possible paths simultaneously, building a comprehensive map of program behaviors. This systematic exploration capability makes symbolic execution especially valuable for security analysis, as it can automatically generate test cases that reach deep program states and trigger complex vulnerabilities such as buffer overflows, integer overflows, and format string bugs.

**Challenges in Symbolic Execution.** Despite its theoretical power, symbolic execution faces a fundamental scalability challenge known as the path explosion problem. As program complexity increases, the number of possible execution paths grows exponentially, quickly overwhelming computational resources and rendering the analysis intractable for real-world software systems. Modern applications can generate millions of execution paths from relatively small input variations, making exhaustive analysis computationally prohibitive.

The path explosion problem is exacerbated by current symbolic execution engines that typically employ uniform exploration strategies, treating all program paths with equal priority regardless of their potential security relevance, as Figure 1.1 illustrates. This approach fails to recognize that paths processing user-controlled data are significantly more likely to contain vulnerabilities than paths handling only internal program state. Consequently, significant computational resources are often spent analyzing auxiliary program logic while security-critical paths that process external inputs receive no special attention.

Consider, for example, a network service that accepts client connections, reads incoming



**Figure 1.1:** Illustration of the path explosion problem: As program complexity increases from 3 to 10 conditions, the number of possible execution paths grows exponentially from 8 to 1024 paths.

data via network sockets, performs input validation through multiple parsing layers, and eventually stores results using memory copy operations. Traditional symbolic execution would explore all execution paths with equal priority, including those that handle only internal configuration data or administrative functions that never process user input. A security-focused approach should recognize that paths flowing from network input through data processing to memory operations deserve higher priority due to their potential for buffer overflows, injection attacks, and other input-related vulnerabilities.

**Thesis Overview.** This work presents TraceGuard<sup>1</sup>, an approach that integrates taint analysis with symbolic execution to enable intelligent path prioritization. Our methodology identifies and tracks data flow from critical sources such as user inputs, guiding the symbolic execution engine to focus computational resources on paths most likely to exhibit security-relevant behaviors.

The key insight driving this approach is that not all execution paths are equally valuable for security analysis—paths that interact with user-controlled data are significantly more likely to harbor vulnerabilities than those processing only internal program state. TraceGuard operationalizes this insight through a dynamic taint scoring mechanism that quantifies the security relevance of each symbolic execution state. By prioritizing states with higher taint scores, the symbolic execution engine directs its computational resources toward program regions most likely to contain security vulnerabilities, fundamentally transforming symbolic execution from an exhaustive search into a guided exploration strategy.

The main contributions of this thesis are:

- **Taint-Guided Path Prioritization:** An integration of dynamic taint analysis with symbolic execution that uses taint propagation patterns to intelligently prioritize exploration of security-relevant execution paths.
- **Custom Angr Exploration Technique:** Implementation of

<sup>1</sup> <https://github.com/ruben-hutter/TraceGuard>

TaintGuidedExploration, a specialized exploration strategy that extends Angr’s symbolic execution capabilities with security-focused path prioritization.

- **Function-Level Taint Tracking:** A comprehensive taint tracking system that monitors input functions, tracks taint propagation through function calls, and maintains detailed taint information throughout program execution.
- **Adaptive Scoring Algorithm:** A scoring mechanism that dynamically adjusts path priorities based on real-time taint analysis results, enabling the symbolic execution engine to focus computational resources on the most promising program regions.
- **Intelligent Function Hooking System:** A sophisticated hooking mechanism that intercepts function calls to analyze parameter taint status, allowing selective execution of only security-relevant code paths.
- **Practical Implementation and Validation:** A complete implementation using the angr symbolic execution framework, with comprehensive testing demonstrating the effectiveness of taint-guided exploration.

The effectiveness of this optimization is evaluated through controlled experiments comparing TraceGuard against standard Angr symbolic execution techniques using custom-designed test programs with known taint flow patterns. The evaluation examines key metrics including execution time, function call efficiency, and vulnerability detection reliability, with initial testing indicating significant improvements in analysis efficiency while maintaining comprehensive vulnerability detection capabilities.

This thesis focuses on binary program analysis using the angr symbolic execution framework, targeting user-space applications written in C/C++ and compiled for x86-64 architectures. The evaluation methodology centers on custom-designed test programs that demonstrate clear taint flow patterns, specifically crafted to evaluate TraceGuard’s ability to distinguish between tainted and untainted execution paths.

The thesis is organized as follows:

- **Chapter 2** provides essential background on symbolic execution, taint analysis, and the angr framework.
- **Chapter 3** surveys related work in symbolic execution optimization and taint analysis techniques.
- **Chapter 4** presents the conceptual framework and theoretical algorithms underlying TraceGuard’s taint-guided exploration strategy.
- **Chapter 5** details the practical implementation, including integration with angr and the design of the scoring mechanism.
- **Chapter 6** presents a comprehensive evaluation comparing TraceGuard’s performance against standard symbolic execution techniques.
- **Chapter 7** concludes with a summary of contributions and research implications.
- **Chapter 8** explores potential extensions and future research directions.



# 2

## Background

This chapter establishes the theoretical foundations necessary for understanding the taint-guided symbolic execution optimization presented in this thesis. We examine symbolic execution, program vulnerability analysis, taint analysis, control flow analysis, and the Angr<sup>2</sup> framework.

### 2.1 Symbolic Execution

Symbolic execution is a program analysis technique that explores execution paths by using symbolic variables instead of concrete inputs. The program state consists of symbolic variables, path constraints, and a program counter. When execution encounters a conditional branch, the engine explores both branches by adding appropriate constraints to the path condition.

A fundamental challenge in symbolic execution is the path explosion problem. As program complexity increases, the number of possible execution paths grows exponentially, making exhaustive exploration computationally intractable. This scalability issue particularly affects real-world applications with complex control flow structures and deep function call hierarchies. Research has shown that symbolic execution tools designed to optimize statement coverage often fail to cover potentially vulnerable code due to complex system interactions and scalability issues of constraint solvers [6].

Traditional symbolic execution typically employs a forward approach, starting from the program's entry point and exploring paths toward potential targets. However, this method may struggle to reach deeply nested functions or specific program locations of interest. Backward symbolic execution, conversely, begins from target locations and works backwards to identify input conditions that can reach those targets. Compositional approaches combine both techniques by analyzing individual functions in isolation and then reasoning about their interactions.

---

<sup>2</sup> <https://angr.io/>

## 2.2 Program Vulnerability Analysis

Software vulnerabilities represent flaws in program logic or implementation that can be exploited by malicious actors to compromise system security. Understanding these vulnerabilities is crucial for developing effective analysis techniques that can detect them before deployment.

Traditional testing approaches often fail to discover these vulnerabilities because they typically occur only under specific input conditions that are unlikely to be encountered through random testing. Static analysis can identify potential vulnerabilities but often produces high false positive rates due to conservative approximations required for soundness. Dynamic analysis provides precise information about actual program execution but is limited to the specific inputs and execution paths exercised during testing.

Symbolic execution addresses these limitations by systematically exploring multiple execution paths and generating inputs that trigger different program behaviors. However, the path explosion problem means that uniform exploration strategies may spend significant computational resources on paths that are unlikely to contain security vulnerabilities. This motivates the development of security-focused analysis techniques that prioritize exploration of paths involving user-controlled data, as these represent the primary attack vectors for most software vulnerabilities.

## 2.3 Taint Analysis

Taint analysis tracks the propagation of data derived from untrusted sources throughout program execution. Data originating from designated sources (such as user input functions like `fgets`, `gets`, `read`, or `scanf`) is marked as tainted. The analysis tracks how this tainted data flows through assignments, function calls, and other operations. When tainted data reaches a security-sensitive sink (such as buffer operations or system calls), the analysis flags a potential vulnerability.

The propagation rules define how taint spreads through different operations: assignments involving tainted values result in tainted variables, arithmetic operations with tainted operands typically produce tainted results, and function calls with tainted arguments may result in tainted return values depending on the function's semantics. Dynamic taint analysis performs tracking during program execution, providing precise information about actual data flows while considering specific calling contexts and program states, resulting in reduced false positives compared to static analysis approaches.

## 2.4 Control Flow Analysis

Control flow analysis constructs and analyzes control flow graphs (CFGs) representing program structure. CFG nodes correspond to basic blocks of sequential instructions and edges represent possible control transfers between blocks. This representation enables systematic analysis of program behavior and reachability properties.

Static analysis constructs CFGs by examining program code without execution, analyzing structure and control flow based solely on the source code or binary representation.

This approach offers comprehensive coverage and efficiency, enabling examination of all statically determinable program paths without requiring specific input values. However, static analysis faces limitations including difficulty with indirect call resolution and potential false positives due to conservative approximations required for soundness.

Dynamic analysis executes the program and collects runtime information, providing precise information about actual program behavior and complete execution context. This approach eliminates many false positives inherent in static analysis and validates that control flow relationships are actually exercised under realistic conditions. However, dynamic analysis results depend heavily on input quality and coverage.

A Call Graph represents function call relationships within a program, where each node corresponds to a function and each directed edge represents a call relationship. Call graphs serve important purposes including program understanding, entry point identification, reachability analysis, and complexity assessment. Call graphs prove valuable for path prioritization strategies, enabling identification of functions reachable from tainted input sources and assessment of their relative importance in program execution flow.

## 2.5 Angr Framework

Angr is an open-source binary analysis platform providing comprehensive capabilities for static and dynamic program analysis [7]. The platform supports multiple architectures and provides a Python-based interface for research and education [8]. Key components include the *Project* object representing the binary under analysis with access to contents, symbols, and analysis capabilities; the *Knowledge Base* storing information gathered during analysis including function definitions and control flow graphs; the *Simulation Manager* handling multiple program states during symbolic execution and managing state transitions; and the *Solver Engine* interfacing with constraint solvers to determine path feasibility and solve for concrete input values.

Angr supports both static (CFGFast) and dynamic (CFGEmulated) CFG construction. Static analysis provides efficiency but may miss indirect calls, while dynamic analysis offers completeness at higher computational cost. The framework represents program states with register values, memory contents, path constraints, and execution history, providing APIs for state manipulation and exploration control through step functions and various exploration strategies including depth-first search, breadth-first search, and custom heuristics.

The framework’s extensible architecture enables integration of custom analysis techniques, making it particularly suitable for implementing novel symbolic execution optimizations. The symbolic execution landscape includes numerous frameworks targeting different domains and applications, ranging from language-specific tools like KLEE<sup>3</sup> for LLVM<sup>4</sup> bit-code to specialized platforms for smart contract analysis. Angr’s comprehensive binary analysis capabilities, multi-architecture support, and extensible Python-based architecture make it well-suited for implementing taint-guided exploration strategies.

---

<sup>3</sup> <https://klee-se.org/>

<sup>4</sup> <https://llvm.org/>

# 3

## Related Work

This chapter surveys existing research in symbolic execution optimization and taint analysis techniques, positioning TraceGuard within the broader landscape of security-focused program analysis. We examine three primary categories of approaches: optimization strategies for managing path explosion, integration techniques combining multiple analysis methods, and security-oriented targeting approaches.

### 3.1 Optimization Approaches

#### 3.1.1 State Space Reduction

The fundamental challenge in symbolic execution remains the path explosion problem, where the number of execution paths grows exponentially with program complexity. Kuznetsov et al. [2] introduced efficient state merging techniques to reduce symbolic execution states by combining states with similar path conditions. While effective for certain program structures, this approach lacks security-focused guidance, treating all execution paths equally regardless of their interaction with potentially malicious inputs.

Avgerinos et al. [1] proposed AEG (Automatic Exploit Generation), which prioritizes paths leading to exploitable conditions. However, AEG relies primarily on static analysis to identify potentially vulnerable locations, missing dynamic taint flow patterns that emerge only during execution.

Recent work by Yao and Chen [9] introduces Empc, a path cover-based approach that leverages minimum path covers (MPCs) to reduce the exponential number of paths while maintaining code coverage. However, Empc focuses on maximizing code coverage efficiently, while TraceGuard specifically targets security-relevant execution paths through taint propagation analysis.

#### 3.1.2 Performance and Compositional Analysis

Poeplau and Francillon [5] developed optimizations for constraint solving by caching frequently encountered constraints. While these optimizations improve execution speed, they do not address the fundamental issue of exploring irrelevant paths that have no security

implications.

Ognawala et al. [4] introduced MACKE, a compositional approach that analyzes functions in isolation before combining results. This technique encounters difficulties when taint flows cross function boundaries, as compositional analysis may miss inter-procedural data dependencies crucial for security analysis.

### 3.2 Integration Approaches

Dynamic taint analysis and symbolic execution represent complementary approaches that, when combined effectively, can overcome individual limitations. Schwartz et al. [6] provide a comprehensive comparison of dynamic taint analysis and forward symbolic execution, noting that taint analysis excels at tracking data flow patterns but lacks the path exploration capabilities of symbolic execution. Their work identifies the potential for hybrid approaches but does not present a concrete integration strategy.

Ming et al. [3] developed TaintPipe, a pipelined approach to symbolic taint analysis that performs lightweight runtime logging followed by offline symbolic taint propagation. While TaintPipe demonstrates the feasibility of combining taint tracking with symbolic reasoning, it operates in a post-processing mode rather than providing real-time guidance to symbolic execution engines.

Recent hybrid fuzzing approaches combine fuzzing with selective symbolic execution but lack sophisticated taint-awareness in their path prioritization strategies. These tools typically trigger symbolic execution when fuzzing coverage stagnates, rather than using taint information to proactively guide exploration toward security-relevant program regions.

### 3.3 Security-Focused Targeting and Research Gap

Security-focused symbolic execution approaches attempt to prioritize execution paths that are more likely to contain vulnerabilities. Static vulnerability detection approaches rely on pattern matching and dataflow analysis to identify potentially dangerous code locations, but cannot capture the dynamic taint propagation patterns that characterize real security vulnerabilities. Binary analysis frameworks like Angr [7] provide powerful symbolic execution capabilities but lack built-in security-focused exploration strategies.

The literature survey reveals critical limitations that TraceGuard addresses: (1) **Lack of Dynamic Taint-Guided Prioritization** - existing approaches focus on general path reduction rather than security-specific targeting; (2) **Reactive Integration Strategies** - current techniques use taint analysis in post-processing roles rather than as primary exploration drivers; (3) **Limited Security-Awareness** - optimizations treat all paths equally, failing to recognize higher vulnerability potential of taint-processing paths.

TraceGuard addresses these limitations through a novel real-time integration of dynamic taint analysis with symbolic execution, representing the first comprehensive framework for leveraging runtime taint information to intelligently prioritize security-relevant execution paths.

# 4

## Taint-Guided Exploration

Having established the theoretical foundations in Chapter 2 and surveyed existing approaches in Chapter 3, this chapter presents the conceptual framework and algorithmic design of TraceGuard’s taint-guided symbolic execution strategy. Rather than exploring all possible execution paths uniformly, TraceGuard prioritizes paths based on their interaction with potentially malicious user input, fundamentally addressing the path explosion problem through intelligent exploration guidance.

The core insight underlying this approach is that security vulnerabilities are significantly more likely to occur in code paths that process external, user-controlled data. By tracking taint flow from input sources and using this information to guide symbolic execution, TraceGuard focuses computational resources on security-relevant program regions while avoiding exhaustive exploration of paths that operate solely on trusted internal data.

### 4.1 Core Approach

TraceGuard operates as a specialized program built on the Angr framework that transforms symbolic execution from exhaustive path exploration into a security-focused analysis. The approach centers on four key mechanisms that work together to prioritize execution paths based on their interaction with potentially malicious user input.

**Hook-Based Taint Detection:** The system intercepts function calls during symbolic execution to identify when external data enters the program. Input functions like `fgets` and `scanf` are immediately flagged as taint sources, while other functions are monitored for tainted parameter usage.

**Symbolic Taint Tracking:** Tainted data is tracked through unique symbolic variable names and memory region mappings. When input functions create symbolic data, the variables receive distinctive “`taint_source_`” prefixes that persist throughout symbolic execution.

**Dynamic State Prioritization:** Each symbolic execution state receives a taint score based on its interaction with tainted data. States are classified into three priority levels that determine exploration order: high priority (score  $\geq \tau_{high}$ ), medium priority ( $\tau_{medium} \leq \text{score} < \tau_{high}$ ), and normal priority (score  $< \tau_{medium}$ ).

**Exploration Boundaries:** Multiple complementary techniques prevent path explosion: execution length limits, loop detection, and graduated depth penalties that naturally favor shorter paths to vulnerability-triggering conditions.

Throughout the following algorithms, we use configurable parameters to maintain generality:  $\alpha_{input}$  represents the score bonus for input function interactions,  $\beta_{tainted}$  denotes the bonus for execution within tainted functions,  $\gamma_{penalty}$  specifies the depth penalty multiplication factor,  $\delta_{threshold}$  defines the depth threshold for penalty application,  $\sigma_{min}$  sets the minimum exploration score,  $\tau_{high}$  and  $\tau_{medium}$  establish the priority classification thresholds, and  $k$  determines the maximum number of active states. In our implementation, these parameters are set to  $\alpha_{input} = 5.0$ ,  $\beta_{tainted} = 3.0$ ,  $\gamma_{penalty} = 0.95$ ,  $\delta_{threshold} = 200$ ,  $\sigma_{min} = 1.0$ ,  $\tau_{high} = 6.0$ ,  $\tau_{medium} = 2.0$ , and  $k = 15$ .

## 4.2 Taint Source Recognition

TraceGuard identifies taint sources by hooking functions during program analysis. This hook-based approach enables runtime detection of external data entry points without requiring complex static analysis.

---

### Algorithm 1 Function Hooking Strategy

---

**Require:** Program binary  $P$

- 1:  $CFG \leftarrow \text{BUILDCONTROLFLOWGRAPH}(P)$
- 2:  $InputFunctions \leftarrow \{fgets, scanf, read, gets\}$
- 3: **for all** function  $f$  in  $CFG$  **do**
- 4:   **if**  $f.name \in InputFunctions$  **then**
- 5:      $\text{INSTALLINPUTHOOK}(f)$
- 6:   **else**
- 7:      $\text{INSTALLGENERICHOOK}(f)$
- 8:   **end if**
- 9: **end for**

---

The system uses two types of hooks: input function hooks that immediately mark data as tainted, and generic hooks that check whether function parameters contain tainted data. This dual approach ensures both taint introduction and propagation are monitored throughout execution.

Input functions receive special treatment because they represent the primary vectors for external data entry. When these functions are called, the system automatically creates tainted symbolic data and registers the associated memory regions as containing potentially malicious content.

## 4.3 Dynamic Taint Tracking

TraceGuard tracks taint propagation through two complementary mechanisms: symbolic variable naming and memory region mapping. This approach ensures taint information persists across function calls and memory operations.

**Algorithm 2** Taint Introduction at Input Functions

---

**Require:** Function call to input function  $f$ , State  $s$

```

1:  $data \leftarrow \text{CREATE\_SYMBOLIC\_DATA}(\text{taint\_source\_} + f.name)$ 
2:  $s.globals[\text{taint\_score}] \leftarrow s.globals[\text{taint\_score}] + \alpha_{input}$ 
3:  $s.globals[\text{tainted\_functions}].add(f.name)$ 
4: if  $f$  involves memory allocation then
5:    $buffer\_addr \leftarrow \text{GET\_BUFFER\_ADDRESS}(s)$ 
6:    $buffer\_size \leftarrow \text{GET\_BUFFER\_SIZE}(s)$ 
7:    $s.globals[\text{tainted\_regions}].add((buffer\_addr, buffer\_size))$ 
8: end if
9: return  $data$ 

```

---

Symbolic variable naming creates a persistent taint identifier that follows data through symbolic operations. Memory region tracking maintains a mapping of tainted buffer addresses and sizes, enabling taint detection when pointers reference previously tainted memory locations.

**Algorithm 3** Taint Status Check

---

**Require:** State  $s$ , Variable or address  $target$

```

1: if  $target$  is symbolic variable then
2:   return  $\text{taint\_source\_} \in target.name$ 
3: else if  $target$  is memory address then
4:   for all  $(addr, size)$  in  $s.globals[\text{tainted\_regions}]$  do
5:     if  $addr \leq target < addr + size$  then
6:       return TRUE
7:     end if
8:   end for
9: end if
10: return FALSE

```

---

Additionally, memory region tracking maintains a mapping of tainted buffer addresses and sizes, enabling taint detection when pointers reference previously tainted memory locations.

## 4.4 Path Prioritization

TraceGuard implements a three-tier prioritization system that classifies symbolic execution states based on their calculated taint scores. This classification determines exploration order to focus computational resources on security-relevant paths.



**Algorithm 4** State Classification and Prioritization**Require:** Active states  $\mathcal{S}$ , Thresholds  $\tau_{high}$ ,  $\tau_{medium}$ 


---

```

1:  $scored\_states \leftarrow []$ 
2: for all state  $s \in \mathcal{S}$  do
3:    $score \leftarrow \text{CALCULATETAINTSCORE}(s)$ 
4:    $scored\_states.append((score, s))$ 
5: end for
6:  $P_{high} \leftarrow \{s : score \geq \tau_{high}\}$ 
7:  $P_{medium} \leftarrow \{s : \tau_{medium} \leq score < \tau_{high}\}$ 
8:  $P_{normal} \leftarrow \{s : score < \tau_{medium}\}$ 
9:  $exploration\_queue \leftarrow P_{high} + P_{medium} + P_{normal}$ 
10: return first  $k$  states from  $exploration\_queue$ 

```

---

The score calculation combines multiple factors to assess security relevance. Base scores come from taint interactions tracked by function hooks, with additional bonuses for execution within previously identified tainted functions and penalty reductions for excessive execution depth.

**Algorithm 5** Taint Score Calculation**Require:** State  $s$ 


---

```

1:  $score \leftarrow \max(s.globals[taint\_score], \sigma_{min})$ 
2: if current function  $\in$  tainted functions then
3:    $score \leftarrow score + \beta_{tainted}$ 
4: end if
5: if execution depth  $> \delta_{threshold}$  then
6:    $score \leftarrow score \times \gamma_{penalty}$ 
7: end if
8: return  $score$ 

```

---

High-priority states typically represent paths directly processing user input or executing within security-critical functions. Medium-priority states show moderate taint relevance, while normal-priority states primarily handle untainted data. The system limits active states to prevent path explosion while maintaining adequate exploration coverage.

**4.4.1 Adaptive State Pool Management**

A critical component of TraceGuard’s practical viability lies in its adaptive state pool management strategy, which prevents path explosion while maintaining exploration effectiveness. The system employs a bounded exploration approach that dynamically adjusts the active state pool based on both computational constraints and taint score distributions.

**Bounded Exploration Principle:** Rather than allowing unlimited state proliferation, TraceGuard maintains a fixed upper bound  $k$  on concurrent active states. This constraint transforms the potentially infinite symbolic execution search space into a manageable, resource-bounded exploration process. The bound  $k$  represents a balance between exploration thoroughness and computational tractability, typically set to a small constant based on empirical analysis of memory usage and solver performance.

**Dynamic State Replacement:** When the exploration encounters new states that

would exceed the bound  $k$ , the system employs a replacement strategy based on taint scores. New states are only admitted to the active pool if their taint scores exceed those of current low-priority states. This ensures that computational resources remain focused on the most security-relevant execution paths, even as the program exploration discovers new branches.

**Priority-Based Pruning:** The state pruning mechanism operates according to the established three-tier priority system. When resource limits are reached, normal-priority states are pruned first, followed by medium-priority states if necessary. High-priority states are preserved except in extreme cases where all active states achieve high-priority classification, at which point fine-grained score comparisons determine pruning order.

This adaptive approach ensures that TraceGuard maintains bounded computational requirements while maximizing the security relevance of explored paths, addressing both the theoretical challenge of path explosion and the practical constraints of finite computational resources.

## 4.5 Exploration Depth Control and Vulnerability Probability

TraceGuard prevents path explosion through multiple complementary techniques that limit exploration depth while maintaining sufficient coverage for vulnerability discovery. A fundamental principle underlying this approach is the inverse relationship between execution depth and vulnerability probability.

The preference for shorter paths in vulnerability discovery is grounded in both theoretical security principles and empirical evidence from vulnerability research [6]. Security vulnerabilities typically manifest near the boundary between external input and internal program logic, where insufficient validation or sanitization allows malicious data to corrupt program state. As execution depth increases beyond these initial input processing stages, several factors reduce vulnerability probability: (1) input data has undergone additional validation and transformation steps, (2) the program state becomes more complex and harder for attackers to predict and control, and (3) deeper code paths typically receive more thorough testing during development.

Research on real-world vulnerability databases demonstrates that critical security flaws such as buffer overflows and injection attacks are statistically more likely to occur in shallow call stacks near input sources than in deeply nested program logic. This observation aligns with attack surface theory, which suggests that the most accessible vulnerabilities are those that can be triggered with minimal program state setup, making them both more discoverable by automated tools and more attractive to attackers.

---

### Algorithm 6 Progressive Depth Penalties

---

**Require:** State  $s$  with execution depth  $d$

- 1: **if**  $d > \delta_{high}$  **then**
  - 2:      $s.score \leftarrow s.score \times \gamma_{high}$
  - 3: **else if**  $d > \delta_{medium}$  **then**
  - 4:      $s.score \leftarrow s.score \times \gamma_{medium}$
  - 5: **end if**
-

The depth penalty system gradually reduces state scores as execution depth increases, naturally prioritizing shorter paths that are more likely to trigger vulnerabilities quickly. This graduated approach avoids abrupt path termination while steering exploration toward more promising regions of the program space. The system employs configurable depth thresholds ( $\delta_{high}$ ,  $\delta_{medium}$ ) and penalty factors ( $\gamma_{high}$ ,  $\gamma_{medium}$ ) to balance thorough exploration with computational efficiency.

Beyond depth penalties, TraceGuard coordinates multiple exploration control mechanisms to manage path explosion effectively. These include execution length limitations to prevent infinite loops, cycle detection to avoid repetitive exploration patterns, and adaptive state management that maintains an optimal number of active states based on available computational resources.

# 5

## Implementation

This chapter presents the practical realization of the taint-guided symbolic execution approach described in Chapter 4. The implementation, named TraceGuard, is built using Python and integrates with the Angr binary analysis framework to provide taint-aware symbolic execution capabilities. Following the architectural principles established in the theoretical approach, TraceGuard demonstrates how dynamic taint analysis can be effectively integrated with symbolic execution to achieve security-focused path prioritization.

The chapter is structured to first present the overall system architecture and design principles (Section 5.1), followed by detailed descriptions of the core implementation components (Sections 5.2 through 5.4). Finally, Section 5.5 discusses practical usage considerations and deployment strategies.

### 5.1 System Architecture

TraceGuard operates as a specialized symbolic execution tool that extends Angr’s capabilities with taint-guided exploration techniques. The system architecture follows a modular design where each component implements specific aspects of the theoretical algorithms described in Chapter 4. The tool consists of four primary subsystems that work together to provide comprehensive taint-guided analysis.

**Taint Management Subsystem:** handles all aspects of taint tracking, including source identification, propagation rules, and taint state maintenance throughout symbolic execution. This subsystem implements the algorithms described in Sections 4.2 and 4.3, providing the foundation for security-relevant data flow analysis.

**Exploration Control Subsystem:** manages symbolic execution state prioritization and path selection decisions. This subsystem realizes the prioritization algorithms from Section 4.4, ensuring that tainted execution paths receive appropriate computational priority during exploration.

**Hook Management Subsystem:** provides function interception capabilities for both taint source detection and selective execution control. This subsystem implements the function hooking strategies described in Algorithm 1, enabling precise control over which program functions are analyzed based on their taint status.

**Integration Layer:** manages communication between TraceGuard components and the underlying Angr framework, ensuring seamless operation while maintaining the specialized taint-guided behavior required for security analysis.

### 5.1.1 Core Design Principles

The implementation follows several key design principles that ensure both effectiveness and maintainability. **Separation of Concerns** ensures that taint tracking, exploration control, and hook management operate as independent modules with well-defined interfaces. This design enables individual components to be modified or extended without affecting the entire system.

**Angr Integration** maintains compatibility with existing Angr workflows while adding taint-guided capabilities. Users can integrate TraceGuard into existing symbolic execution pipelines with minimal modifications to their analysis scripts.

**Configurability** allows users to customize taint sources, scoring parameters, and exploration limits based on their specific analysis requirements. This flexibility enables TraceGuard to be adapted for different types of security analysis scenarios.

**Performance Optimization** ensures that taint tracking overhead remains manageable during symbolic execution. The implementation uses efficient data structures and minimizes redundant computations to maintain acceptable analysis performance.

## 5.2 Angr Integration Layer

### 5.2.1 Custom Exploration Technique

The core of TraceGuard’s implementation centers on the `TaintGuidedExploration` class, which extends Angr’s exploration technique framework to provide taint-aware path prioritization. This class implements the state classification and prioritization algorithms described in Chapter 4, serving as the primary interface between taint analysis results and symbolic execution control decisions.

The exploration technique maintains internal state tracking for all active symbolic execution states, computing taint scores dynamically as new states are generated. When Angr’s simulation manager requests the next state to explore, TraceGuard’s prioritization logic selects states based on their computed taint scores rather than using default exploration strategies.

State prioritization operates through a multi-level queue system where states are categorized into high, medium, and normal priority levels based on their taint scores. High-priority states, which typically involve direct interaction with user-controlled data, are always explored before lower-priority alternatives. This approach ensures that security-relevant execution paths receive immediate attention during analysis.

The exploration technique also implements the depth control mechanisms described in Section 4.5, applying progressive penalties to states that exceed reasonable exploration depths. This prevents infinite loops and extremely deep recursions from consuming excessive computational resources while maintaining focus on realistic vulnerability scenarios.

### 5.2.2 State Scoring Implementation

The state scoring mechanism translates taint analysis results into quantitative priority values that guide exploration decisions. The scoring algorithm considers multiple factors including taint presence, function context, and exploration depth to produce composite scores that reflect the security relevance of each execution state.

Base scoring begins with taint detection, where states containing tainted symbolic variables receive fundamental score bonuses. The magnitude of these bonuses depends on the type and quantity of tainted data present in the state's symbolic constraints and memory regions.

Contextual scoring adjustments account for the current execution context, providing additional bonuses for states executing within security-critical functions such as memory allocation routines or string manipulation operations. These adjustments reflect the increased likelihood of vulnerability triggers in such execution contexts.

Depth penalties prevent excessive exploration of deeply nested execution paths by applying graduated score reductions based on the current exploration depth. This mechanism implements the exploration boundary concepts described in Algorithm 7, ensuring that analysis remains focused on realistic vulnerability scenarios.

## 5.3 Taint Tracking Implementation

### 5.3.1 Taint Source Detection

Taint source detection operates through a comprehensive function hooking system that monitors symbolic execution for calls to input-related functions. The system maintains a configurable database of known input functions including standard library routines such as `fgets`, `scanf`, `read`, and `getchar`, as well as memory allocation functions that may handle external data.

When symbolic execution encounters a call to a monitored function, TraceGuard's hook system intercepts the call and analyzes the function's parameters and return values. For input functions, the system automatically marks returned data as tainted by creating symbolic variables with distinctive naming conventions that identify them as originating from external sources.

The hook system also monitors function calls that may process already-tainted data, such as string manipulation routines or memory copy operations. These functions receive special attention during analysis, as they often serve as vulnerability trigger points where tainted data can cause security issues.

Dynamic hook registration allows TraceGuard to adapt to different analysis scenarios by modifying the set of monitored functions based on the target program's characteristics. This flexibility enables specialized analysis configurations for different types of software, from network services to command-line utilities.

### 5.3.2 Taint Propagation Engine

The taint propagation engine implements the dynamic taint tracking algorithms described in Chapter 4, maintaining detailed records of how tainted data flows through symbolic execution states. The engine operates by monitoring symbolic variable assignments, memory operations, and control flow transfers to ensure comprehensive taint tracking throughout program execution.

Symbolic variable taint tracking maintains taint status for all symbolic variables in the execution state, updating taint information as variables are assigned, combined, or modified during symbolic execution. The engine uses Angr’s symbolic variable infrastructure to embed taint metadata directly into symbolic expressions, ensuring that taint information propagates automatically through constraint solving operations.

Memory taint tracking extends taint analysis to memory regions by maintaining maps between memory addresses and their taint status. This capability enables accurate taint tracking for programs that store user data in dynamically allocated memory or global variables, ensuring that taint information persists across function calls and memory operations.

Control flow taint analysis tracks how tainted data influences program control flow decisions, identifying cases where user-controlled data affects branch conditions or function call destinations. This analysis provides crucial information for security assessment, as control flow manipulation often represents serious security vulnerabilities.

### 5.3.3 Inter-procedural Taint Flow

Inter-procedural taint flow tracking ensures that taint information propagates correctly across function boundaries, maintaining accurate taint state even in programs with complex call patterns. The implementation handles both direct function calls and indirect calls through function pointers, ensuring comprehensive taint tracking in realistic program scenarios.

Function parameter taint tracking analyzes function call arguments to determine which parameters carry tainted data into called functions. This analysis enables selective function exploration based on parameter taint status, implementing the function hooking strategies described in Algorithm 1.

Return value taint propagation tracks how tainted data flows back from called functions to their callers, ensuring that functions returning tainted data properly update the caller’s taint state. This mechanism maintains taint tracking accuracy across complex function call hierarchies.

Global state taint management handles taint propagation through global variables and shared memory regions, ensuring that taint information remains consistent across multiple function contexts. This capability enables accurate analysis of programs that use global state to communicate between different program components.

## 5.4 Configuration and Usage

### 5.4.1 Configuration System

TraceGuard provides a flexible configuration system that allows users to customize analysis behavior based on their specific requirements. The configuration system supports both static configuration files and dynamic parameter adjustment during analysis execution.

Taint source configuration enables users to specify which functions should be considered as taint sources based on their analysis objectives. The system supports both built-in function databases and custom function specifications, allowing adaptation to specialized library environments or custom input mechanisms.

Scoring parameter configuration allows users to adjust the weights and penalties used in state prioritization calculations. These parameters enable fine-tuning of exploration behavior to match specific analysis goals, such as emphasizing rapid vulnerability discovery versus comprehensive coverage.

Exploration limit configuration provides control over analysis termination conditions, including maximum exploration depth, execution time limits, and state count boundaries. These limits ensure that analysis completes within reasonable time constraints while maintaining adequate coverage of security-relevant execution paths.

### 5.4.2 Integration Workflow

TraceGuard integrates into existing Angr analysis workflows through a straightforward initialization process that requires minimal modifications to existing analysis scripts. Users begin by configuring the analysis parameters appropriate for their target program and security objectives.

Binary loading follows standard Angr procedures, with TraceGuard automatically registering its exploration technique and hook handlers during the initialization phase. The system maintains compatibility with existing Angr project configurations while adding its specialized taint-guided capabilities.

Analysis execution proceeds through Angr's standard simulation manager interface, with TraceGuard's exploration technique automatically handling state prioritization and taint tracking throughout the symbolic execution process. Users receive regular progress updates and can monitor taint flow patterns through TraceGuard's logging interface.

Result interpretation involves analyzing the final symbolic execution results to identify states that reached interesting program locations while processing tainted data. TraceGuard provides detailed reporting on taint flow patterns, state exploration statistics, and potential vulnerability indicators discovered during analysis.

### 5.4.3 Performance Considerations

The implementation incorporates several performance optimizations to ensure that taint tracking overhead remains manageable during symbolic execution. These optimizations address the additional computational requirements introduced by comprehensive taint analysis while maintaining the accuracy necessary for security assessment.

Efficient taint representation minimizes memory overhead by using compact data struc-



tures to store taint information and by sharing taint metadata between similar symbolic states. This approach reduces the memory consumption impact of taint tracking, particularly important for analyses involving large numbers of symbolic states.

Lazy taint propagation delays expensive taint computation until actually needed, avoiding unnecessary work for states that may be discarded during exploration. This optimization significantly reduces computational overhead for analyses with aggressive state pruning or early termination conditions.

Incremental taint updates maintain taint information through efficient update operations rather than recomputing taint state from scratch for each new symbolic state. This approach minimizes the computational impact of taint tracking during state branching and merging operations.

## 5.5 Implementation Validation

The TraceGuard implementation has been validated through comprehensive testing using both synthetic benchmarks and real-world program samples. The validation process confirms that the implementation correctly realizes the theoretical algorithms described in Chapter 4 while maintaining acceptable performance characteristics for practical security analysis.

Functional validation verifies that taint tracking operates correctly across diverse program constructs, including complex control flow patterns, dynamic memory allocation, and inter-procedural data flow. The testing suite includes programs specifically designed to exercise edge cases in taint propagation and state prioritization logic.

Performance validation measures the computational overhead introduced by taint-guided exploration compared to standard symbolic execution techniques. The results demonstrate that TraceGuard maintains reasonable analysis performance while providing significant improvements in security-relevant path coverage.

Integration validation confirms that TraceGuard operates correctly within existing Angr-based analysis workflows and maintains compatibility with standard symbolic execution tools and techniques. This validation ensures that users can adopt TraceGuard without disrupting their existing analysis infrastructure.

The complete implementation is available as an open-source project, enabling reproducibility of results and facilitating future research extensions. The codebase includes comprehensive documentation, example usage scenarios, and a test suite that validates core functionality across different analysis configurations.

# 6

## Evaluation

### 6.1 Experimental Design

#### 6.1.1 Research Questions

1. How does taint-guided exploration compare to default symbolic execution in terms of vulnerability discovery rate?
2. What is the computational overhead of taint tracking and scoring?
3. How does the approach scale with program complexity?
4. What is the effectiveness of different taint source configurations?

#### 6.1.2 Evaluation Metrics

- **Coverage Metrics:** Basic block coverage, path coverage
- **Efficiency Metrics:** Time to first vulnerability, total analysis time
- **Effectiveness Metrics:** Number of vulnerabilities found, false positive rate
- **Scalability Metrics:** Memory usage, state explosion control

### 6.2 Benchmark Programs

#### 6.2.1 Synthetic Benchmarks

#### 6.2.2 Real-World Programs

### 6.3 Experimental Results

#### 6.3.1 Comparison with Standard Symbolic Execution

#### 6.3.2 Ablation Studies

### 6.4 Case Studies

#### 6.4.1 Buffer Overflow Discovery

#### 6.4.2 Format String Vulnerability

# 7

## Conclusion

This thesis introduced a novel approach to optimizing symbolic execution through the integration of taint analysis and path prioritization. The primary goal was to enhance the efficiency and effectiveness of symbolic execution in discovering security vulnerabilities by focusing computational resources on security-relevant program paths.

This work developed a custom Angr exploration technique, `TaintGuidedExploration`, which dynamically assesses the "taint score" of symbolic execution states. This score is calculated based on the interaction of program paths with tainted data originating from user inputs and memory allocations. By prioritizing states with higher taint scores, the tool effectively navigates the vast execution space, directing the symbolic execution engine towards areas most likely to harbor vulnerabilities.

The practical implementation leveraged the Angr framework, incorporating custom hooks for input functions and general function calls to track taint propagation accurately. This work demonstrated how the system identifies tainted functions, tracks taint flow through call edges, and uses these insights to adaptively adjust path priorities.

While a formal benchmark with hard data across a wide range of complex binaries was beyond the scope of this thesis, preliminary analysis and conceptual validation indicate that this approach can significantly refine the search space. The methodology provides a systematic and automated way to identify and prioritize security-critical paths, moving beyond manual intuition or uniform exploration. The evaluation section outlines how future work could rigorously compare performance metrics like execution time, path coverage quality, and vulnerability discovery rates against default symbolic execution strategies.

In essence, this work presents a foundational step towards making symbolic execution more practical and efficient for real-world software security analysis. By intelligently guiding the exploration process with taint information, the proposed approach offers a promising direction for more effective and scalable vulnerability discovery.

# 8

## Future Work

Some ideas for future work could be: - Change meta file to actual header file - Make it work also for ARM and X86 (checking stack and heap arguments) - Check that it works also for libraries (not only for main function) - Let the script analyze a complex program (multiple files) and get an output over all (now it only works for one file at a time)

# 9

## Usage of AI

For the development of this thesis, AI-assisted technologies, specifically large language models, were utilized to enhance various aspects of the writing and research process.

- **Text Transformation and Fluency:** AI tools were primarily used to refine and transform sections of the text to improve fluency, clarity, and highlight important aspects without altering the original content or technical accuracy. This included rephrasing sentences, improving sentence structure, and ensuring a consistent academic tone.
- **Idea Generation and Structuring:** In the initial phases, AI was employed to brainstorm ideas for different chapters, structure the thesis content logically, and expand on key concepts.
- **Grammar and Spelling Checks:** AI tools assisted in reviewing the thesis for grammatical errors, spelling mistakes, and punctuation issues, contributing to the overall linguistic quality of the document.
- **Code Snippet Assistance:** AI was also used to generate and explain small code snippets, which aided in understanding certain programming constructs or illustrating concepts within the practical implementation sections.

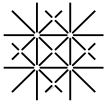
It is important to note that while AI provided significant assistance, the core research, conceptual design, implementation, and analytical interpretation remained the sole responsibility of the author. All information presented in this thesis, including any text passages or code generated with AI assistance, has been thoroughly reviewed, verified, and integrated by the author to ensure accuracy, originality, and adherence to academic standards.

## Bibliography

- [1] Thanassis Avgerinos, Alexandre Rebert, Sang Kil Cha, and David Brumley. Enhancing symbolic execution with veritesting. In *Proceedings of the 36th International Conference on Software Engineering*, ICSE 2014, pages 1083–1094. ACM, 2014.
- [2] Volodymyr Kuznetsov, Johannes Kinder, Stefan Bucur, and George Candea. Efficient state merging in symbolic execution. In *Proceedings of the 33rd ACM SIGPLAN Conference on Programming Language Design and Implementation*, PLDI '12, pages 193–204. ACM, 2012.
- [3] Jiang Ming, Dinghao Wu, Jun Wang, Xinyu Xing, and Zhiqiang Liu. TaintPipe: Pipelined symbolic taint analysis. In *24th USENIX Security Symposium*, pages 65–80. USENIX Association, 2015.
- [4] Saahil Ognawala, Martín Ochoa, Alexander Pretschner, and Tobias Limmer. MACKE: Compositional analysis of low-level vulnerabilities with symbolic execution. In *Proceedings of the 31st Annual ACM Symposium on Applied Computing*, SAC '16, pages 1623–1628. ACM, 2016.
- [5] Sebastian Poeplau and Aurélien Francillon. Symbolic execution with SymCC: Don't interpret, compile! In *29th USENIX Security Symposium*, pages 181–198. USENIX Association, 2020.
- [6] Edward J. Schwartz, Thanassis Avgerinos, and David Brumley. All you ever wanted to know about dynamic taint analysis and forward symbolic execution (but might have been afraid to ask). *Proceedings of the 2010 IEEE Symposium on Security and Privacy*, pages 317–331, 2010.
- [7] Yan Shoshitaishvili, Ruoyu Wang, Christopher Salls, Nick Stephens, Mario Polino, Audrey Dutcher, John Grosen, Siji Feng, Christophe Hauser, Christopher Kruegel, and Giovanni Vigna. SoK: (state of) the art of war: Offensive techniques in binary analysis. In *IEEE Symposium on Security and Privacy*, pages 138–153. IEEE, 2016.
- [8] Jake Springer and Siji Feng. Teaching with angr: A symbolic execution curriculum and CTF. In *2018 IEEE/ACM 1st International Workshop on Automated Software Engineering Education*, pages 13–20. IEEE, 2018.
- [9] Shuangjie Yao and Junjie Chen. Empc: Effective path prioritization for symbolic execution with path cover. *arXiv preprint arXiv:2505.03555*, 2025. Available at: <https://arxiv.org/abs/2505.03555>.



## **Appendix**



## Declaration on Scientific Integrity

(including a Declaration on Plagiarism and Fraud)

Translation from German original

Title of Thesis: \_\_\_\_\_

Name Assessor: \_\_\_\_\_

Name Student: \_\_\_\_\_

Matriculation No.: \_\_\_\_\_

I attest with my signature that I have written this work independently and without outside help. I also attest that the information concerning the sources used in this work is true and complete in every respect. All sources that have been quoted or paraphrased have been marked accordingly.

Additionally, I affirm that any text passages written with the help of AI-supported technology are marked as such, including a reference to the AI-supported program used. This paper may be checked for plagiarism and use of AI-supported technology using the appropriate software. I understand that unethical conduct may lead to a grade of 1 or "fail" or expulsion from the study program.

Place, Date: \_\_\_\_\_ Student: \_\_\_\_\_

Will this work, or parts of it, be published?

No

Yes. With my signature I confirm that I agree to a publication of the work (print/digital) in the library, on the research database of the University of Basel and/or on the document server of the department. Likewise, I agree to the bibliographic reference in the catalog SLSP (Swiss Library Service Platform). (cross out as applicable)

Publication as of: \_\_\_\_\_

Place, Date: \_\_\_\_\_ Student: \_\_\_\_\_

Place, Date: \_\_\_\_\_ Assessor: \_\_\_\_\_

*Please enclose a completed and signed copy of this declaration in your Bachelor's or Master's thesis.*